

```

(A) $ $HIVE_HOME/bin hive --service cli
(B) hive> set hive.cli.print.current.db=true;
(C) hive (default)> CREATE DATABASE ourfirstdatabase; OK
Time taken: 3.756 seconds
(D) hive (default)> USE ourfirstdatabase; OK
Time taken: 0.039 seconds
(E) hive (ourfirstdatabase)> CREATE TABLE our_first_table
(
> FirstName STRING,
> LastName STRING,
> EmployeeId INT);
OK
Time taken: 0.043 seconds
hive (ourfirstdatabase)> quit;

```

1)

```
hduser@ubuntu:~$ start-all.sh
```

```
hduser@ubuntu:~$ hive
```

```
hive> set hive.cli.print.current.db=true;
```

```
hive (default)> create database ourfirstdatabase;
```

```
OK
```

```
Time taken: 1.955 seconds
```

```
hive (default)> use ourfirstdatabase;
```

```
OK
```

```
Time taken: 0.048 seconds
```

```
hive (ourfirstdatabase)> create table our_first_table(firstname string,lastname string,employeeid int);
```

```
OK
```

```
Time taken: 0.873 seconds
```

```
hive> show tables;
```

```
OK
```

```
our_first_table
```

```
Time taken: 0.114 seconds, Fetched: 1 row(s)
```

```
//set db properties use alter command
```

```
hive> ALTER DATABASE ourfirstdatabase SET DBPROPERTIES
```

```
> ('creator'=Swapnali Ware',  
> 'created_for'='Learning Hive DDL');
```

OK

Time taken: 0.315 seconds

```
hive> DESCRIBE DATABASE EXTENDED ourfirstdatabase;
```

OK

```
ourfirstdatabase      hdfs://localhost:54310/user/hive/warehouse/ourfirstdatabase.db  
      hduser USER {created_for=Learning Hive DDL, creator=Swapnali Ware}
```

Time taken: 0.091 seconds, Fetched: 1 row(s)

```
hive> DROP DATABASE ourfirstdatabase CASCADE;
```

OK

Time taken: 1.691 seconds

2)

Table name Customer Information

Rows	Column Families				
	ContactInfo		CustomerName		
	EA	SA	FN	MN	LN
1	sbw.ware@sinhgad.edu	Lonavala	Swapnali	B	Ware
2	riya.singh@xyz.com	Pune	Riya	S	Singh

```
hbase(main):005:0> create 'CustomerInformation','ContactInfo','CustomerName'
```

0 row(s) in 2.5900 seconds

```
=> Hbase::Table - CustomerInformation
```

```
hbase(main):006:0> list
```

TABLE

CustomerInformation

1 row(s) in 0.1070 seconds

```
=> ["CustomerInformation"]
```

```
hbase(main):001:0> list
```

TABLE

CustomerInformation

1 row(s) in 0.5790 seconds

(main):002:0> put 'CustomerInformation',1,'ContactInfo:EA','sbw.sit@sinhgad.edu'

0 row(s) in 0.6280 seconds

hbase(main):003:0> put 'CustomerInformation',1,'ContactInfo:SA','Lanavala'

0 row(s) in 0.0200 seconds

hbase(main):004:0> put 'CustomerInformation',1,'CustomerName:FN','Swapnali'

0 row(s) in 0.0530 seconds

hbase(main):005:0> put 'CustomerInformation',1,'CustomerName:MN','B'

0 row(s) in 0.0290 seconds

hbase(main):006:0> put 'CustomerInformation',1,'CustomerName:LN','Ware'

0 row(s) in 0.0040 seconds

hbase(main):009:0> put 'CustomerInformation',2,'ContactInfo:EA','riya.singh@xyz.com'

0 row(s) in 0.0060 seconds

hbase(main):010:0> put 'CustomerInformation',2,'ContactInfo:SA','Pune'

0 row(s) in 0.0140 seconds

hbase(main):011:0> put 'CustomerInformation',2,'CustomerName:FN','Riya'

0 row(s) in 0.0090 seconds

hbase(main):012:0> put 'CustomerInformation',2,'CustomerName:MN','S'

0 row(s) in 0.0170 seconds

hbase(main):013:0> put 'CustomerInformation',2,'CustomerName:LN','Singh'

0 row(s) in 0.0160 seconds

hbase(main):014:0> scan 'CustomerInformation'

ROW	COLUMN+CELL
-----	-------------

1	column=ContactInfo:EA, timestamp=1512450536862, value=sbw.
---	--

sit@sinhgad.edu

```
1      column=ContactInfo:SA, timestamp=1512450573649, value=Lana
      vala
1      column=CustomerName:FN, timestamp=1512450617894, value=Sw
      pnali
1      column=CustomerName:LN, timestamp=1512450643726, value=War
      e
1      column=CustomerName:MN, timestamp=1512450630986, value=B
2      column=ContactInfo:EA, timestamp=1512450750624, value=riya
      .singh@xyz.com
2      column=ContactInfo:SA, timestamp=1512450771492, value=Pune
2      column=CustomerName:FN, timestamp=1512450784697, value=Riy
      a
2      column=CustomerName:LN, timestamp=1512450840361, value=Sin
      gh
2      column=CustomerName:MN, timestamp=1512450819560, value=S
```

2 row(s) in 0.2040 seconds

Goto Hive terminal

- In Step (A), you create an external table with a Key field to link up with the HBase row keys (1 and 2), and two **map data types (name and info)** to link up with the two column families (ContactInfo and CustomerName).
- Note the syntax for providing this linkage via the **WITH SERDEPROPERTIES** keywords. This SerDe configuration technique is quite common in Hive DDL.
- Note as well that the **TBLPROPERTIES** keyword is crucial for connecting the new **external hive_hbase_table** with the actual CustomerInformation HBase table name.
- Step (B) shows how the key value pairs in HBase ({"FN","Swapnali"}, for example) are now available for querying with the help of the HiveQL. Note the syntax for accessing the Hive map data type in Step (C). You can select the value of the info map type using the notation info ["EA"] where "EA" is the key.

```
hive> CREATE EXTERNAL TABLE hive_hbase_table (
      > key INT,
```

```
> name map<STRING,STRING>,  
> info map<STRING, STRING>)  
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'  
> WITH SERDEPROPERTIES ("hbase.columns.mapping" =  
> ": key, CustomerName:, ContactInfo:")  
> TBLPROPERTIES ("hbase.table.name" = "CustomerInformation");
```

OK

Time taken: 6.248 seconds

```
hive> SELECT * FROM hive_hbase_table;
```

OK

```
1      {"FN":"Swapnali","LN":"Ware","MN":"B"}  
      {"EA":"sbw.sit@sinhgad.edu","SA":"Lanavala"}  
2      {"FN":"Riya","LN":"Singh","MN":"S"}  
      {"EA":"riya.singh@xyz.com","SA":"Pune"}
```

Time taken: 2.343 seconds, Fetched: 2 row(s)

```
hive> SELECT info["EA"] FROM hive_hbase_table WHERE  
      > name["FN"] = "Swapnali" AND name["LN"] = "Ware";
```

OK

```
CREATE TABLE IF NOT EXISTS FlightInfo2007
```

```
(
```

```
Year SMALLINT, Month TINYINT, DayofMonth TINYINT,
```

```
DayOfWeek TINYINT,
```

```
DepTime SMALLINT, CRSDepTime SMALLINT, ArrTime SMALLINT,CRSArrTime SMALLINT,
```

```
UniqueCarrier STRING, FlightNum STRING, TailNum STRING,
```

```
ActualElapsedTime SMALLINT, CRSElapsedTime SMALLINT,
```

```

AirTime SMALLINT, ArrDelay SMALLINT, DepDelay SMALLINT,
Origin STRING, Dest STRING, Distance INT,
TaxiIn SMALLINT, TaxiOut SMALLINT, Cancelled SMALLINT,
CancellationCode STRING, Diverted SMALLINT,
CarrierDelay SMALLINT, WeatherDelay SMALLINT,
NASDelay SMALLINT, SecurityDelay SMALLINT,
LateAircraftDelay
SMALLINT)
COMMENT 'Flight InfoTable'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE
TBLPROPERTIES ('creator'='PSB ', 'created_at'='Tues Dec 5 3:00:00 EDT 2017');
OK
> LINES TERMINATED BY '\n'

```

Time taken: 0.292 seconds

```
hive> load data local inpath '/home/student/Desktop/2007.csv' into table FlightInfo2007;
```

```
hive> CREATE TABLE IF NOT EXISTS FlightInfo2008 LIKE FlightInfo2007;
```

```
hive> load data local inpath '/home/hduser/Desktop/2008.csv' into table FlightInfo2008;
```

```

hive> CREATE TABLE IF NOT EXISTS myFlightInfo (
    Year SMALLINT, DontQueryMonth TINYINT, DayofMonth
    TINYINT, DayOfWeek TINYINT, DepTime SMALLINT, ArrTime SMALLINT,
    UniqueCarrier STRING, FlightNum STRING,
    AirTime SMALLINT, ArrDelay SMALLINT, DepDelay SMALLINT,
    Origin STRING, Dest STRING, Cancelled SMALLINT,
    CancellationCode STRING)
COMMENT 'Flight InfoTable'

```

```
PARTITIONED BY(Month TINYINT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LINES TERMINATED BY '\n'
STORED AS RCFILE TBLPROPERTIES ('creator'='swapnali ware',
'created_at'='Mon sep 2 14:24:19 EDT 2017');
```

OK

Time taken: 1.697 seconds

```
hive> INSERT OVERWRITE TABLE myflightinfo PARTITION (Month=1)
```

```
> SELECT Year, Month, DayofMonth, DayOfWeek, DepTime,
> ArrTime, UniqueCarrier,FlightNum, AirTime, ArrDelay, DepDelay, Origin,
> Dest, Cancelled,CancellationCode FROM FlightInfo2008 WHERE Month=1;
```

```
hive> FROM FlightInfo2008 INSERT INTO TABLE myflightinfo
```

```
> PARTITION (Month=2) SELECT Year, Month, DayofMonth, DayOfWeek, DepTime,
> ArrTime, UniqueCarrier, FlightNum,
> AirTime, ArrDelay, DepDelay, Origin, Dest, Cancelled,
> CancellationCode WHERE Month=2
> INSERT INTO TABLE myflightinfo
> PARTITION (Month=12)
> SELECT Year, Month, DayofMonth, DayOfWeek, DepTime,
> ArrTime, UniqueCarrier, FlightNum,
> AirTime, ArrDelay, DepDelay, Origin, Dest, Cancelled,
> CancellationCode WHERE Month=12;
```

```
hive> SHOW PARTITIONS myflightinfo;
```

OK

month=1

month=12

month=2

Time taken: 0.344 seconds, Fetched: 3 row(s)

hive> CREATE TABLE myflightinfo2007 AS

```
> SELECT Year, Month, DepTime, ArrTime, FlightNum,  
> Origin, Dest FROM FlightInfo2007  
> WHERE (Month = 7 AND DayofMonth = 3) AND  
> (Origin='JFK' AND Dest='ORD');
```

hive>SELECT * FROM myFlightInfo2007;

OK

2007	7	700	834	5447	JFK	ORD
2007	7	1633	1812	5469	JFK	ORD
2007	7	1905	2100	5492	JFK	ORD
2007	7	1453	1624	4133	JFK	ORD
2007	7	1810	1956	4392	JFK	ORD
2007	7	643	759	903	JFK	ORD
2007	7	939	1108	907	JFK	ORD
2007	7	1313	1436	915	JFK	ORD
2007	7	1617	1755	917	JFK	ORD
2007	7	2002	2139	919	JFK	ORD

Time taken: 1.219 seconds, Fetched: 10 row(s)

hive> CREATE TABLE myFlightInfo2008 AS

```
> SELECT Year, Month, DepTime, ArrTime, FlightNum,  
> Origin, Dest FROM FlightInfo2008  
> WHERE (Month = 7 AND DayofMonth = 3) AND  
> (Origin='JFK' AND Dest='ORD');
```

hive> SELECT * FROM myFlightInfo2008;

OK

2008	7	930	1103	5199	JFK	ORD
2008	7	705	849	5687	JFK	ORD
2008	7	1645	1914	5469	JFK	ORD

2008	7	1345	1514	4392	JFK	ORD
2008	7	1718	1907	1217	JFK	ORD
2008	7	757	929	1323	JFK	ORD
2008	7	928	1057	907	JFK	ORD
2008	7	1358	1532	915	JFK	ORD
2008	7	1646	1846	917	JFK	ORD
2008	7	2129	2341	919	JFK	ORD

Time taken: 0.424 seconds, Fetched: 10 row(s)

JOIN

Hive>SELECT m8.Year, m8.Month, m8.FlightNum, m8.Origin, m8.Dest, m7.Year, m7.Month, m7.FlightNum, m7.Origin, m7.Dest FROM myFlightinfo2008 m8 JOIN myFlightinfo2007 m7 ON m8.FlightNum=m7.FlightNum;

2008	7	5469	JFK	ORD	2007	7	5469	JFK	ORD
2008	7	4392	JFK	ORD	2007	7	4392	JFK	ORD
2008	7	907	JFK	ORD	2007	7	907	JFK	ORD
2008	7	915	JFK	ORD	2007	7	915	JFK	ORD
2008	7	917	JFK	ORD	2007	7	917	JFK	ORD
2008	7	919	JFK	ORD	2007	7	919	JFK	ORD

hive> SELECT m8.FlightNum,m8.Origin,m8.Dest,m7.FlightNum,m7.Origin,m7.Dest FROM myFlightinfo2008 m8 FULL OUTER JOIN myFlightinfo2007 m7 ON m8.FlightNum=m7.FlightNum;

1217	JFK	ORD	NULL	NULL	NULL
1323	JFK	ORD	NULL	NULL	NULL
NULL	NULL	NULL	4133	JFK	ORD
4392	JFK	ORD	4392	JFK	ORD
5199	JFK	ORD	NULL	NULL	NULL
NULL	NULL	NULL	5447	JFK	ORD
5469	JFK	ORD	5469	JFK	ORD
NULL	NULL	NULL	5492	JFK	ORD
5687	JFK	ORD	NULL	NULL	NULL
NULL	NULL	NULL	903	JFK	ORD

907	JFK	ORD	907	JFK	ORD
915	JFK	ORD	915	JFK	ORD
917	JFK	ORD	917	JFK	ORD
919	JFK	ORD	919	JFK	ORD

Time taken: 10.33 seconds, Fetched: 14 row(s)

```
hive>SELECT
m8.Year,m8.Month,m8.FlightNum,m8.Origin,m8.Dest,m7.Year,m7.Month,m7.FlightNum,m7.Origin,m7
.Dest FROM myFlightinfo2008 m8 LEFT OUTER JOIN myFlightinfo2007 m7 ON
m8.FlightNum=m7.FlightNum;
```

2008	7	5199	JFK	ORD	NULL	NULL	NULL	NULL	NULL
2008	7	5687	JFK	ORD	NULL	NULL	NULL	NULL	NULL
2008	7	5469	JFK	ORD	2007	7	5469	JFK	ORD
2008	7	4392	JFK	ORD	2007	7	4392	JFK	ORD
2008	7	1217	JFK	ORD	NULL	NULL	NULL	NULL	NULL
2008	7	1323	JFK	ORD	NULL	NULL	NULL	NULL	NULL
2008	7	907	JFK	ORD	2007	7	907	JFK	ORD
2008	7	915	JFK	ORD	2007	7	915	JFK	ORD
2008	7	917	JFK	ORD	2007	7	917	JFK	ORD
2008	7	919	JFK	ORD	2007	7	919	JFK	ORD

```
hive> CREATE INDEX f08_index ON TABLE flightinfo2008 (Origin) AS
```

```
> 'COMPACT' WITH DEFERRED REBUILD;
```

OK

Time taken: 1.124 seconds

```
hive> ALTER INDEX f08_index ON flightinfo2008 REBUILD;
```

```
hive>SHOW INDEXES ON FlightInfo2008;
```

OK

f08_index	flightinfo2008	origin	default__flightinfo2008_f08_index__
compact			

Time taken: 2.549 seconds, Fetched: 1 row(s)

```
hive> SELECT Origin, COUNT(1) FROM
```

```
> flightinfo2008 WHERE Origin = 'SYR' GROUP BY Origin;
```

```
hive> DESCRIBE default__flightinfo2008_f08_index__;
```

OK

origin	string
--------	--------

_bucketname	string
-------------	--------

_offsets	array<bigint>
----------	---------------

Time taken: 0.927 seconds, Fetched: 3 row(s)

```
hive> SELECT Origin, SIZE(`_offsets`)
```

```
> FROM default__flightinfo2008_f08_index__ WHERE origin='SYR';
```

OK

SYR	12032
-----	-------

Time taken: 0.705 seconds, Fetched: 1 row(s)

```
hive> CREATE VIEW avgdepdelay AS
```

```
> SELECT DayOfWeek, AVG(DepDelay) FROM
```

```
> FlightInfo2008 GROUP BY DayOfWeek;
```

```
hive> SELECT * FROM avgdepdelay;
```

3	8.289761053658728
---	-------------------

6	8.645680904903614
---	-------------------

1	10.269990244459473
---	--------------------

4	9.772897177836702
---	-------------------

7	11.568973392595312
---	--------------------

2	8.97689712068735
---	------------------

5	12.158036387869656
----------	---------------------------

Day 5 under the results in Step (B) — had the highest number of delays.

Step (A): We want to point out that Hive's Data Definition Language (DDL) also includes the `CREATE VIEW` statement, which can be quite useful. In Hive, views allow a query to be saved but data is not stored as with the Create Table as Select (CTAS) statement.

When a view is referenced in HiveQL, Hive executes the query and then uses the results which could be part of a larger query. This can be very useful to simplify complex queries and break them down into logical components.

Additionally, the `GROUP BY` clause, which gathers all the days per week and allows the `AVG` aggregate function to provide a consolidated answer per day.

After we answered our question above about average flight delays per day,